

Week 07 1

CSC209 Fall 2023

Dr. Demetres (dee) Kostas

October 26, 2023

Announcements

- Reading week next week
- A3 is up
- office hours after this lecture

Processes

- what is a process?
- a data-structure
 - a particular abstraction
- that represents an **individual**
 - program that can run
 - on a system

As a datastructure

- variables needed to execute
 - the actual machine instructions
- as well as indicators for
 - where you currently are
- and also the memory model
 - of the current program

Running processes

- processes ultimately
 - feature machine instructions
 - that need to run on HW
- But there are too many
 - so they share the HW
- The OS facilitates this

What does this mean?

- essentially, that
 - at worst
 - you don't know when a process
 - might be interrupted
 - * (often after system calls)
- It will continue again
 - exactly where it left off
- but order is unknown

Process IDs (PID)

- most coordination of processes
 - is handled automatically
 - by the OS
- but PID allows us
 - as developers
 - * and by proxy, our apps
 - to monitor other processes

Processes in a linux system

- suppose you boot linux
- this means the computer boots some code
 - called the kernel
 - which prepares the filesystem
 - and enables the system calls
- Then, it creates PID 1
 - a.k.a init
 - which sits around to adopt orphans

After init

- init is somehow duplicated
 - PID 2 is just another copy of init
 - with the same variables
 - * with it's own memory locations for
 - each of the variables
- this new copy
 - get's cannabalized into the next
 - * step in booting the system
- and this model underlies
 - all process creation

Consider if PID 2 was a shell

- what if the exact next program
- was a shell program
 - connected to screen and keyboard
- we could use it to start more!
 - other programs!
 - or maybe a window manager
 - * to contain more programs

Fork

- the first system call
- that we need to learn about
- to be process masters

What is fork?

- essentially it duplicates
 - the current process' datastructure
 - * memory model, and current code position
 - creating a **child**
 - as if there was a big struct (PCB)
 - somewhere and you just copied it
- this includes
 - the memory model for this process
 - * **not the literal locations**

Why is this important

- because all of the variables
 - are duplicated
 - by they are in different places!
- but *some things*
 - like file descriptors/names
 - or other identifiers
 - * e.g. PID
 - still access the same entities
 - * because they're an ID

So it's the same program?

- not anymore
- once `fork` is executed
- the variables are the same
 - but they are **independent**
 - they just have the same values
 - now there is a place in
 - * literal memory, that is different
 - * for the original and copies
- but the memory model for both is unchanged

Why would the child be any different?

- We use the return value of `fork`!
- child processes get a value of 0
 - while parents get the PID of the child
- if non-zero, this is the parent
- if zero, this is the child

Other system calls

- there is an ecosystem
 - of system calls
- for process management

`wait(pid)`

- without a pid
 - it just waits for *children*
 - of the current process
 - * to completely finish

- with a pid
 - it can retrieve info
 - most importantly
 - * the **exit status**

exec

- this is the crucial companion
 - to **fork**
- while **wait** allows us to
 - make sure our processes
 - * work together as intended
- **exec** allows us to run
 - **new** programs
- not just duplicates

exec

- this is ultimately how
 - an executable
 - gets converted into a process
- give **exec** an executable
 - as an argument
- and it modifies the current process
 - to execute the executable's values

Roughly speaking

- `exec` fills in the process datastructure
 - with new values
 - that it picks up from an executable file
- If the file is a proper
 - executable type (and perm.)
- the values
 - in the current model's addresses
 - are replaced by the executable's values
 - * e.g. code, RO-strings, etc.

These are actually system call families

- there are many system calls
- of the `wait` or `moreso`, `exec` variety
- so what is different
 - not much, besides changes in argument
 - or argument usage
 - the fundamentals are the same

man pages

Let's start an example

WORKSHEET

`fork.pdf` and `fork_fruits.c`