

Week 08 1

CSC209 Fall 2023

Dr. Demetres (dee) Kostas

October 31, 2023

Announcements

- reading week is definitely
 - **next week**
- you have all the concepts
 - that you need to finish A3
 - It is due **Nov. 15th**

Recap

- processes are
 - abstractions that run code
 - * on a system
 - with their own memory model
- **fork** allows us
 - to duplicate an existing process
 - and continue execution in both

Recap cont'd

- `exec` provides
 - a way to swap the contents
 - of one process' internals with
 - * an executable file('s)
 - * i.e. the process becomes
 - a different program
- `wait`
 - simply waits for *child*
 - * processes
 - but required for coordination

More wait

- last class didn't feature
 - too much with `wait`
 - though you saw it in lab

Wait and return codes

- note that `wait(int *status)`
 - waits for **any** child to finish
 - * called within the parent
 - so if you have 3 children
 - * you could call this 3 times
- specifically calling
 - `waitpid(pid_t childpid, int *status, int options);`
 - * allows you to wait
 - for specific children

Status codes for wait

- these are integrated into an int
 - we use macros to access
 - * certain **bit fields**
 - you'll learn more in a few weeks
- for now, you can use:
 - `WIFEXITED(status)`
 - * true if exits normally
 - `WEXITSTATUS(status)`
 - * for the numeric code
 - returned by child's main

Let's do a worksheet together

wait.pdf

Communication between processes

- if processes are separated
 - even if they are duplicates
- how could we ever get them
 - to work together?
- what have we learned so far
 - that might enable this?

Consider using a file

- what if we opened a specific file
- and we write to it
 - with Process A
 - * (and close it nicely)

- and then read from it
 - with Process B
 - * (without interfering with A)

Inter-process communication

- this would allow Process A
 - to write arbitrary bytes
 - later consumed by Process B
- we could have an additional file
 - for the other direction!
 - or share just one file. . .
 - * (this gets more difficult)

System conveniences

- rather than do this
 - and waste disk space
 - and worry about opening
 - * at the right time
 - * without interference
- we have an additional system tool
 - called **pipes**
 - that are exactly what we describe

first recall file-descriptors

- this is extensively discussed
 - in the PCRS videos
- but essentially
 - all the open *file-like* constructs
 - * are assigned a number

- * a file descriptor
- this allows the system to keep track
 - of what each process opens

standard fd's

- this includes
 - stdin: 0
 - stdout: 1
 - stderr: 2
- *as well as pipes!*
- **Note:** files are slightly different
 - the FILE struct is used
 - but there is an underlying
 - * `fileno()`

Piping

- as you might expect
 - we use another system call:
 - * `int pipe(int fd[2]);`
 - the return value
 - * is for error detection
- otherwise two fd's are written
 - to the argument

I don't see how this connects to inter-process communication

- fd's are interesting
 - because `fork` duplicates them
 - so that child processes
 - * **access the same ones**

- i.e. the fd numbers refer to the same
- * file-like abstraction
- provided by the system

A parent-child pipe example

```
#include <unistd.h>

int pipefds[2];
pipe(pipefds); // no error check

int read_end = pipefds[0];
int write_end = pipefds[1];

pid_t result = fork();
if (result == 0){ //child
    close(read_end);
    write(write_end, ..., ...);
    close(write_end);
} else { //parent
    close(write_end);
    read(read_end, ..., ...);
    close(read_end);
}
```

Two more system calls

- `size_t read(int fd, void *dest_buffer, size_t num_to_read);`
 - much like `fread`, but with fds
 - * not with `FILE` structs
- `size_t write(int fd, void *src_buffer, size_t num_to_write)`
 - much like `fwrite`, but with fds