

Week 10 2
CSC209 Fall 2023

Dr. Demetres (dee) Kostas

November 23, 2023

Contents

1	Announcements	1
2	Sockets	1
2.1	What does <code>socket</code> do?	2
2.2	(Continued)	2
2.3	So we have a first system call	2
2.4	<code>int socket(int domain, int type, int protocol)</code>	3
2.5	So... Internet now?	3
3	The “server” end	4
3.1	<code>bind()</code>	4
3.2	<code>listen()</code>	5
4	The “client” end	5
5	Architecture lifecycle	6
6	System calls in detail	6
6.1	We saw <code>socket</code>	7
6.2	<code>int bind(int sockfd, struct sockaddr *addr, socklen_t addrlen)</code>	7
6.3	Code snippet from Lab 10	7
6.4	<code>htons(port_num)</code>	8
6.5	<code>int listen(int sockfd, int backlog)</code>	8
6.6	See Lab 10 for helper functions	9
6.7	<code>int accept()</code>	9
6.8	<code>accept()</code> continued	9
6.9	Connection lifecycle	10

6.10	<code>int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen)</code>	10
6.11	<code>close(sockfd)</code>	10
7	Final remarks	11
7.1	TCP/IP	11
7.2	“Peers”	12
7.3	PF_INET vs AF_INET	12
7.4	sending/receiving data	12

1 Announcements

- Lab 10 this week
 - Lab 11 will be the last!
- After the lecture
 - you should be able to
 - understand what
 - * *needs to be done*
 - for A4

2 Sockets

- recalling pipes
 - when we said `pipe()`
 - * the system created
 - some sort of buffer
 - * and two `fd`’s
 - one to put data
 - one to get data

2.1 What does socket do?

- create something similar. . .
- but gives us 1 `fd`!

- there are various implementations
 - but the socket
 - is essentially an endpoint
 - * (the **fd**)
 - that we can read/write data
 - * **both** on same **fd**

2.2 (Continued)

- this endpoint automatically
 - formats data going in
 - and converts data before reading
 - * according to the args:
 - **domain**
 - IPv4 == **AF_INET**
 - **type**
 - TCP == **SOCK_STREAM**
 - **protocol**
 - 0 for CSC209

2.3 So we have a first system call

- one for creating an end point
 - for both read and write
- represented by a single **fd**
- such that *the system*
 - will basically *package*
 - * the data we write
 - or *unwrap* the data
 - * that comes in for read
- according to the format

2.4 `int socket(int domain, int type, int protocol)`

- return value is the fd
 - or -1 on error
- We mentioned two “packets” (type)
 - TCP and UDP
 - * these are the types
 - `SOCK_STREAM` (TCP)
 - guaranteed error detection
 - `SOCK_DGRAM` (UDP)
 - no guarantees

2.5 So... Internet now?

- not quite
- we’ve made a socket
 - like making a pipe
- let’s connect it to our network

3 The “server” end

- this is the end
- that does not *initiate*
 - the “conversation”
 - “passive” end
- it works by waiting
 - for connections

3.1 bind()

- once we have a socket...
 - “structure”
- we can manually assign it
 - to a specific port (and address)
 - * to wait for connection
- basically, the socket’s `fd`
 - needs to be connected
 - to a specific port
 - * on the network device
 - * so *others* can remote access

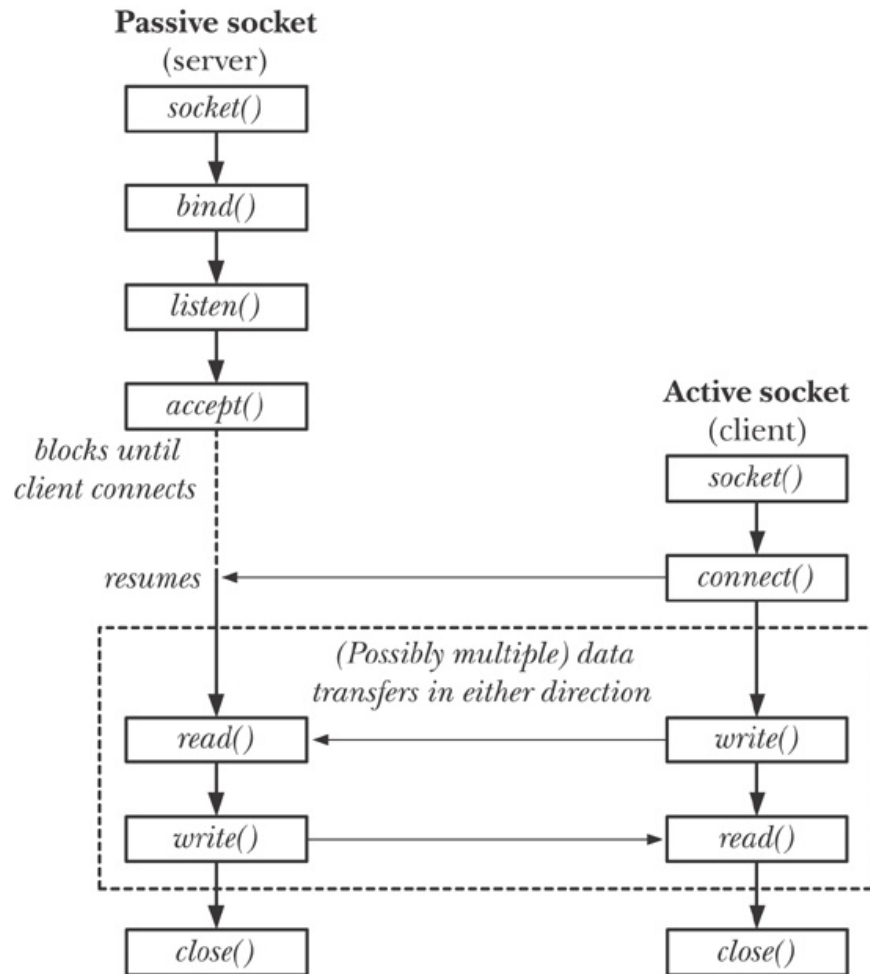
3.2 listen()

- alternatively just say this
- `bind` forces a *specific* port
 - but we could just tell the system
 - that this socket
 - * should *listen* for data
- similar to `bind`
 - but the system just assigns
 - * the socket to a port on the
 - * network interface randomly

4 The “client” end

- this end *requires*
 - a “server” to be available
 - and **connects** to it
- it typically
 - just uses whatever free port
 - * is available
 - it doesn’t affect the user

5 Architecture lifecycle



6 System calls in detail

- let's look at the *additional*
 - system calls that one socket
 - * needs to connect to another
- we understand the architecture

- we just need to use
 - * the available calls correctly

- again the focus is TCP/IPv4

6.1 We saw socket

- 9 times out of 10
 - it will look like this

```
int socket_fd = socket(AF_INET, SOCK_STREAM, 0);
```

6.2 int bind(int sockfd, sockaddr *addr, socklen_t addrlen)

- -1 on error, and sockfd we know
- socklen_t is basically an int
 - its a typedef for holding integers
- addrlen is just the number
 - of bytes used for addr
 - * to **not** overflow

```
struct sockaddr_in {
    sa_family_t    sin_family;    /* AF_INET */
    in_port_t      sin_port;      /* Port number */
    struct in_addr  sin_addr;      /* IPv4 address */
};
```

6.3 Code snippet from Lab 10

```
struct sockaddr_in *init_server_addr(int port) {
    struct sockaddr_in *addr = malloc(sizeof(struct sockaddr_in));

    // Allow sockets across machines.
    addr->sin_family = AF_INET;
```



```

// The port the process will listen on.
addr->sin_port = htons(port);

// Clear this field; sin_zero is used for padding for the struct.
memset(&(addr->sin_zero), 0, 8);

// Listen on all network interfaces.
addr->sin_addr.s_addr = INADDR_ANY;

return addr;
}

```

6.4 htons(port_num)

- remember when we talked about
 - byte order
- int are 4 bytes
 - but left-right, or right-left?
- htons
 - ensures the 16-bit port number
 - * isn't accidentally lost

6.5 int listen(int sockfd, int backlog)

- 0 on success, -1 error
- listen's job is to enable
 - active/client connections
 - but it doesn't initiate
 - * actual data transfer
 - this is for accept()
- As such, backlog is for
 - how many pending connections
 - the system must manage

6.6 See Lab 10 for helper functions

```
struct sockaddr_in *init_server_addr(int port);  
int set_up_server_socket(struct sockaddr_in *self, int num_queue);
```

6.7 int accept()

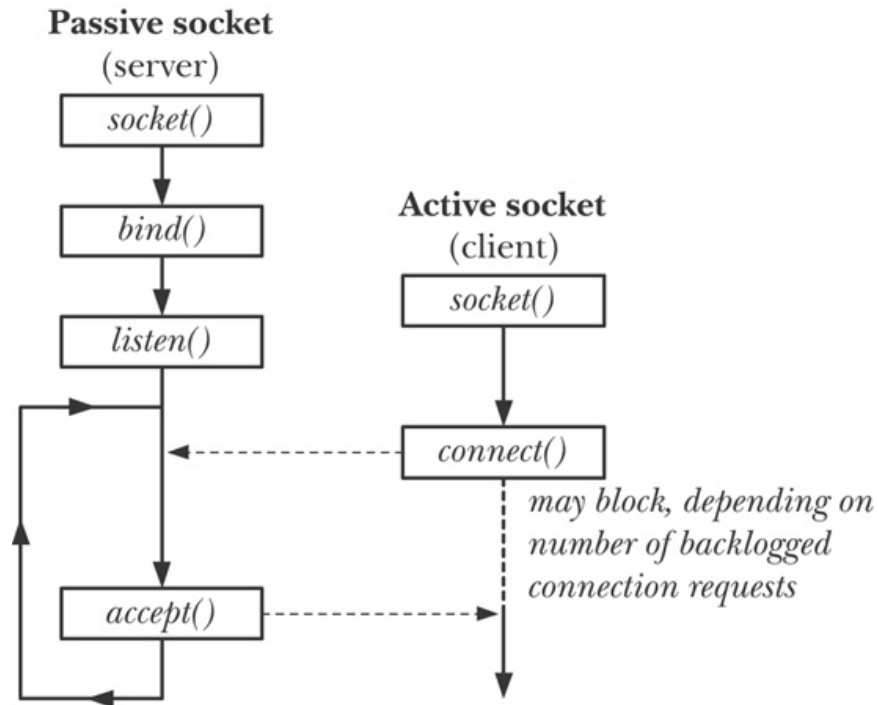
- blocks until there is connection
- creates **another** socket
 - the first one is to handle
 - * connections
- *this one* is to read/write data
- will return a new fd for R/W
 - on success, or -1 on error

```
int rw_fd = accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

6.8 accept() continued

- addr
 - is *filled* by accept!
- what about *addrlen?
 - sizeof(addr)
 - so the system doesn't overflow
- this holds the info
 - about the **active** peer
 - * (client)
- NULL and 0 are valid
 - if we don't care

6.9 Connection lifecycle



6.10 `int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen)`

- same as `bind` really
- but for the “active” side
 - client

- 0 on success, -1 error

6.11 `close(sockfd)`

- if either side closes
 - system cleans up socket resources
- if *the other side* reads

- as in, they haven't closed
 - * they will receive EOF
- if *the other side* write
 - a SIGPIPE signal is raised

7 Final remarks

- a passive process on one machine
 - a.k.a *a server*
- plus any number of active
 - connectors
 - a.k.a. *clients*
- make up, what you might call
 - a client-server application
 - * that uses TCP/IP

7.1 TCP/IP

- take the networking course
 - if you are finding this interesting
- some brief notes
 - every un-corrupted packet
 - * is acknowledge upon recv
 - the stream of bytes are broken
 - * into a sequence of packets
 - * but users don't need to
 - manage this in any way

7.2 “Peers”

- rather than client/server
- I think the terminology
 - of passive vs active peers
 - makes more sense
 - * peer applications
 - on either side of INET socket
- allowing a much more flexible
 - relationship
 - it *is* bidirectional...

7.3 PF_INET vs AF_INET

- assignment uses PF
- they are interchangeable on IPv4

7.4 sending/receiving data

- thus far we’ve seen
 - how to use `send()` and `recv()`
- you could also have used
 - more *raw* approaches
- `read()` and `write()`
 - could have worked
- and there is an additional pair
 - `sendto()` and `recvfrom()`